

# FraudDetection

May 2, 2018

## 1 Problem Description : Classification of Fraudulent Transactions

```
In [1]: from keras.models import Sequential
        from keras.layers import Dense
        from keras.layers import LSTM
        from keras.layers import Dropout
        from keras.layers import Flatten, Input
        from keras import backend as K
        from keras.models import Model, load_model
        from sklearn.preprocessing import StandardScaler
        import matplotlib.pyplot as plt
        from sklearn.model_selection import KFold
        from scipy.spatial import distance
        from sklearn.decomposition import PCA
        from numpy import linalg as LA
        from keras.objectives import categorical_crossentropy
        from sklearn.metrics import roc_curve, auc
        import math
        from scipy.stats import pearsonr
        import copy
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import mean_squared_error, r2_score
        import itertools
        import csv
        from sklearn import metrics
        import tensorflow as tf
        import tensorflow.contrib.layers as tl
        import numpy as np
        import pandas as pd
        from sklearn import linear_model
        from sklearn.ensemble import RandomForestClassifier
        import seaborn as sb
        %matplotlib inline
```

```
/home/ramchalamkr/.local/lib/python2.7/site-packages/h5py/__init__.py:36: FutureWarning: Conver
    from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

## 2 Part 1 : Supervised Learning Approaches

### 3 Steps

#### 3.1 1. Checking the Data

```
In [5]: X = pd.read_csv('fraud_prep.csv', delimiter=',')
        print X.shape
        X.head()
```

(284807, 31)

```
Out [5]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	

	V8	V9	...	V21	V22	V23	V24	\
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	

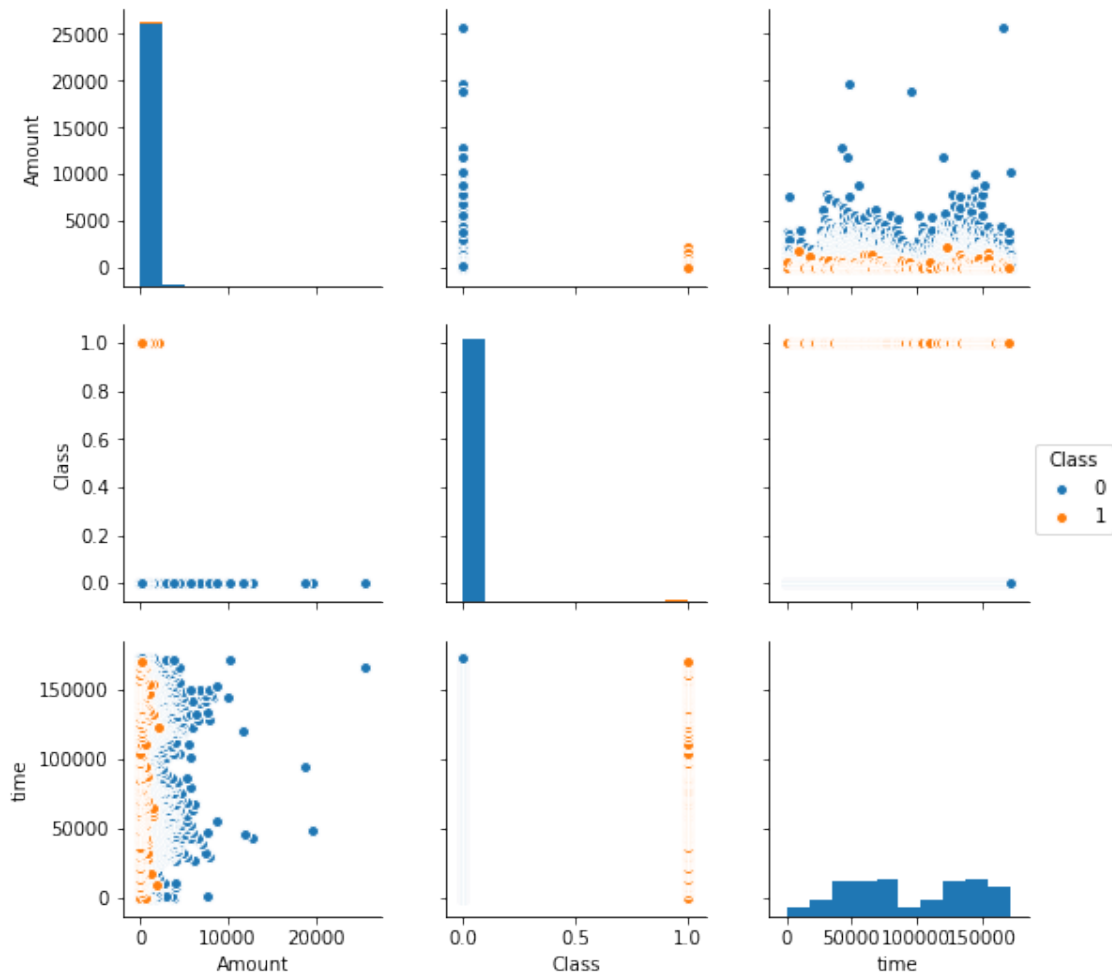
  

	V25	V26	V27	V28	Amount	Class
0	0.128539	-0.189115	0.133558	-0.021053	149.62	0
1	0.167170	0.125895	-0.008983	0.014724	2.69	0
2	-0.327642	-0.139097	-0.055353	-0.059752	378.66	0
3	0.647376	-0.221929	0.062723	0.061458	123.50	0
4	-0.206010	0.502292	0.219422	0.215153	69.99	0

[5 rows x 31 columns]

```
In [3]: temp = pd.DataFrame({'Amount':X['Amount'],'time':X['Time'], 'Class':X['Class']})
        sb.pairplot(temp.dropna(), hue='Class')
```

```
Out [3]: <seaborn.axisgrid.PairGrid at 0x7fda7509dcd0>
```



### 3.1.1 Supervised Methods

```
In [7]: Y = X['Class']
del X['Class']
print X.shape
print Y.shape
```

```
(284807, 30)
(284807,)
```

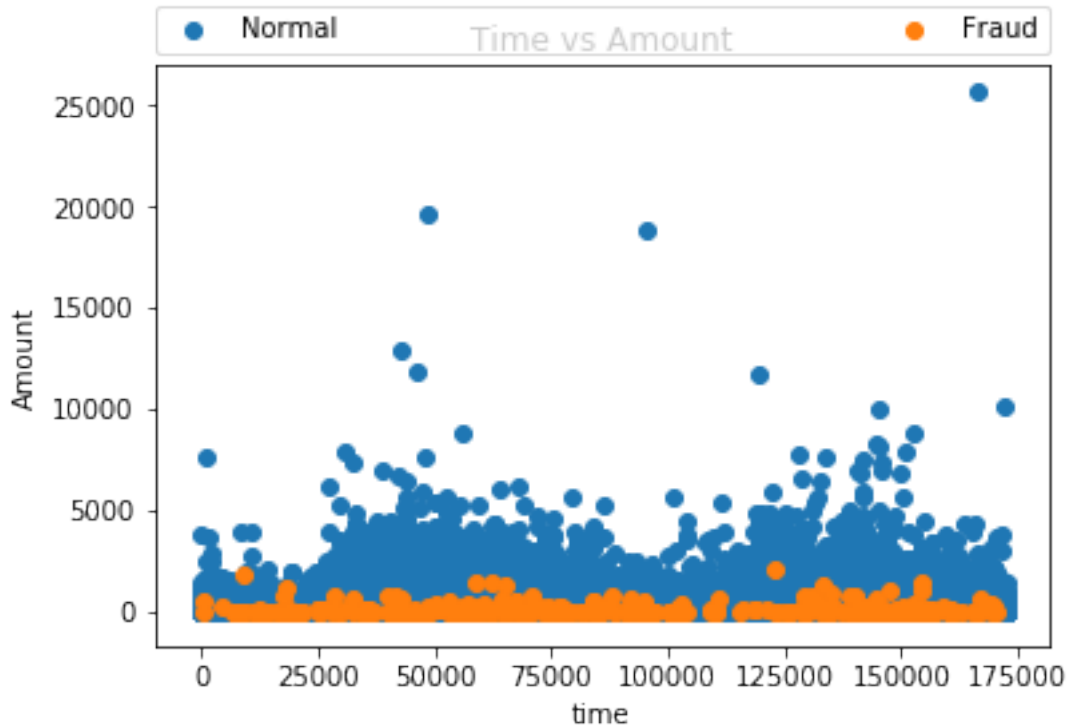
As can be seen, across time, not much distinction in data, but mostly in terms of amount. So can drop time

```
In [4]: time = pd.DataFrame({'Amount':X['Amount'],'time':X['Time'], 'class':Y})
Normaltime = time[(time['class'] == 0)]
FraudTran = time[(time['class'] == 1)]
```

```

plt.scatter(Normaltime['time'], Normaltime['Amount'],label = "Normal")
plt.scatter(FraudTran['time'],FraudTran['Amount'],label = "Fraud")
plt.title("Time vs Amount")
plt.legend(bbox_to_anchor=(0., 1.02, 1., .102), loc=3,
           ncol=2, mode="expand", borderaxespad=0.)
plt.xlabel("time")
plt.ylabel("Amount")
plt.show()

```



## 3.2 2. Tidying the Data

```

In [5]: #Removing mean and scaling to unit variance
        X['Amount'] = StandardScaler().fit_transform(X['Amount'].values.reshape(-1, 1))
        #removing the time column
        del X['Time']

```

### 3.2.1 Logistic Regression (Baseline Model)

```

In [6]: X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.3,random_state=42)
        clf = linear_model.LogisticRegression()
        clf.fit(X_train,Y_train)
        Y_pred = clf.predict(X_test)
        probs = clf.predict_proba(X_test)

```

```

print "Accuracy",metrics.accuracy_score(Y_test, Y_pred)
print "AUC Value", metrics.roc_auc_score(Y_test, probs[:, 1])
print "Confusion Matrix"
print metrics.confusion_matrix(Y_test, Y_pred)
print "Precision Recall"
print metrics.classification_report(Y_test, Y_pred)

```

Accuracy 0.9992626663389628

AUC Value 0.9806636963106952

Confusion Matrix

```

[[85295   12]
 [   51   85]]

```

Precision Recall

	precision	recall	f1-score	support
0	1.00	1.00	1.00	85307
1	0.88	0.62	0.73	136
avg / total	1.00	1.00	1.00	85443

```

In [7]: X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.3,random_state=42)
        clf = RandomForestClassifier(max_depth=20)
        clf.fit(X_train,Y_train)
        Y_pred = clf.predict(X_test)
        probs = clf.predict_proba(X_test)
        print "Accuracy",metrics.accuracy_score(Y_test, Y_pred)
        print "AUC Value", metrics.roc_auc_score(Y_test, probs[:, 1])
        print "Confusion Matrix"
        print metrics.confusion_matrix(Y_test, Y_pred)
        print "Precision Recall"
        print metrics.classification_report(Y_test, Y_pred)

```

Accuracy 0.9995669627705019

AUC Value 0.9526959807449773

Confusion Matrix

```

[[85299    8]
 [   29  107]]

```

Precision Recall

	precision	recall	f1-score	support
0	1.00	1.00	1.00	85307
1	0.93	0.79	0.85	136
avg / total	1.00	1.00	1.00	85443

### 3.2.2 PCA

```
In [8]: print X.shape
        X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.2,random_state=42)

        Covariance = np.dot(X_train.T,X_train)

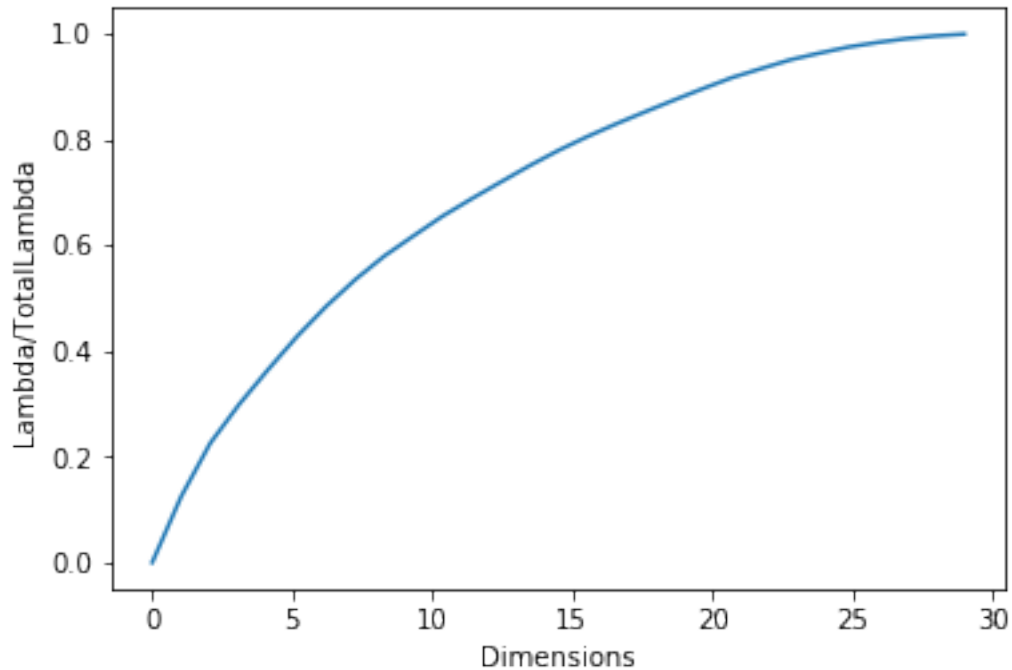
        Lambda, e = LA.eigh(Covariance)
        Lambda = Lambda.reshape(Lambda.shape[0],1)
        Lambda = sorted(Lambda,reverse=True)
        TotalLambda = np.sum(Lambda)
        LambdaProp = []
        for i in range(X_train.shape[1]):
            temp = np.sum(Lambda[0:i])*1.0/TotalLambda
            LambdaProp.append(temp)

        Dim = np.linspace(0, X_train.shape[1], X_train.shape[1])

        plt.plot(Dim,LambdaProp)
        plt.legend(bbox_to_anchor=(0., 1.02, 1., .102), loc=3,
                   ncol=2, mode="expand", borderaxespad=0.)
        plt.xlabel('Dimensions')
        plt.ylabel('Lambda/TotalLambda')
        plt.show()

(284807, 29)

/home/ramchalamkr/.local/lib/python2.7/site-packages/matplotlib/axes/_axes.py:545: UserWarning
  warnings.warn("No labelled objects found. ")
```



3.2.3 This means there is no point in doing PCA as the total variance is represented by mostly all 28~30 features. It could also mean that the data already contains the best features and all are possibly required to get best results.

### 3.3 3. Model Development and Performance

#### 3.3.1 Feed Forward Neural Networks

3.3.2 (Model Architecture followed an extension of a published paper.<https://pdfs.semanticscholar.org/0419/c275f05841d87ab9a4c9767a4f997b61a50e.pdf>)

#### 3.3.3 using original data

```
In [41]: Y = np.reshape(Y,[Y.shape[0],1])
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.3,random_state=42)
X_train = np.asarray(X_train)
X_test = np.asarray(X_test)
Y_train = np.asarray(Y_train)
Y_test = np.asarray(Y_test)
K.clear_session()
InputWidth = X_train.shape[1]
model = Sequential()
model.add(Dense(256, input_shape = (InputWidth,), activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.2))
```

```

model.add(Dense(64, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(32, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(1, activation='sigmoid'))
print model.summary()
# Compile model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
op = model.fit(X_train, Y_train, validation_data=(X_test, Y_test), epochs=10, batch_s
#output = model.fit(X_train, Y_train, validation_data=(X_test, Y_test), epochs=150, b
y_pred = model.predict(X_test)
y_pred = [int(item) for sublist in y_pred for item in sublist]
#print y_pred.shape
print "Accuracy",metrics.accuracy_score(Y_test, y_pred)
#print np.mean(y_pred==Y_test)
print metrics.confusion_matrix(Y_test, y_pred)
print metrics.classification_report(Y_test, y_pred)

```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	7680
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 64)	8256
dropout_3 (Dropout)	(None, 64)	0
dense_4 (Dense)	(None, 32)	2080
dropout_4 (Dropout)	(None, 32)	0
dense_5 (Dense)	(None, 1)	33

```

=====
Total params: 50,945
Trainable params: 50,945
Non-trainable params: 0
=====
None
Train on 199364 samples, validate on 85443 samples
Epoch 1/10
10s - loss: 0.0103 - acc: 0.9979 - val_loss: 0.0037 - val_acc: 0.9994
Epoch 2/10

```



```

11s - loss: 0.0041 - acc: 0.9994 - val_loss: 0.0033 - val_acc: 0.9994
Epoch 3/10
10s - loss: 0.0038 - acc: 0.9994 - val_loss: 0.0030 - val_acc: 0.9994
Epoch 4/10
10s - loss: 0.0035 - acc: 0.9994 - val_loss: 0.0032 - val_acc: 0.9994
Epoch 5/10
11s - loss: 0.0032 - acc: 0.9994 - val_loss: 0.0028 - val_acc: 0.9994
Epoch 6/10
11s - loss: 0.0032 - acc: 0.9993 - val_loss: 0.0028 - val_acc: 0.9994
Epoch 7/10
11s - loss: 0.0030 - acc: 0.9994 - val_loss: 0.0027 - val_acc: 0.9994
Epoch 8/10
14s - loss: 0.0028 - acc: 0.9994 - val_loss: 0.0028 - val_acc: 0.9994
Epoch 9/10
12s - loss: 0.0027 - acc: 0.9994 - val_loss: 0.0028 - val_acc: 0.9994
Epoch 10/10
12s - loss: 0.0028 - acc: 0.9994 - val_loss: 0.0029 - val_acc: 0.9994
0.0
[[85307      0]
 [  128      8]]
      precision    recall  f1-score   support

     0       1.00      1.00      1.00     85307
     1       1.00      0.06      0.11       136

avg / total       1.00      1.00      1.00    85443

```

/home/ramchalamkr/.local/lib/python2.7/site-packages/ipykernel\_launcher.py:27: DeprecationWarning

### 3.3.4 DownSampling for imbalanced class

```

In [47]: #Downsampling the data to handle imbalance class.
X = pd.read_csv('fraud_prep.csv',delimiter=',')
fraud = X[X.Class ==1]
normal = X[X.Class==0]
#fraud.head()
frames = [fraud, normal[0:1000]]

DownSample = pd.concat(frames)
DownSample.head()
Y = DownSample['Class']
del DownSample['Class']
print DownSample.shape
print Y.shape
del DownSample['Time']

```

```

        #Removing mean and scaling to unit variance
        DownSample['Amount'] = StandardScaler().fit_transform(DownSample['Amount'].values.reshape(
(1492, 30)
(1492,)

```

## Only keras

```

In [48]: Y = np.reshape(Y,[Y.shape[0],1])
        class_weight = {0: 1.,1: 50.}
        X_train,X_test,Y_train,Y_test = train_test_split(DownSample,Y,test_size = 0.3,random_s
        X_train = np.asarray(X_train)
        X_test = np.asarray(X_test)
        Y_train = np.asarray(Y_train)
        Y_test = np.asarray(Y_test)
        K.clear_session()
        InputWidth = X_train.shape[1]
        model = Sequential()
        model.add(Dense(256, input_shape = (InputWidth,), activation='relu'))
        model.add(Dropout(0.2))
        model.add(Dense(128, activation='relu'))
        model.add(Dropout(0.2))
        model.add(Dense(64, activation='relu'))
        model.add(Dropout(0.2))
        model.add(Dense(32, activation='relu'))
        model.add(Dropout(0.2))
        model.add(Dense(1, activation='sigmoid'))
        print model.summary()
        # Compile model
        model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
        #output = model.fit(X_train, Y_train, validation_data=(X_test, Y_test), epochs=150, bat
        output = model.fit(X_train, Y_train, validation_data=(X_test, Y_test), epochs=50, bat
        y_pred = model.predict(X_test)
        y_pred = [int(item) for sublist in y_pred for item in sublist]
        #print y_pred.shape
        print "accuracy", metrics.accuracy_score(Y_test, y_pred)
        print metrics.confusion_matrix(Y_test, y_pred)
        print metrics.classification_report(Y_test, y_pred)
        print "predicting on whole dataset"
        Y = X['Class']
        del X['Class']
        print X.shape
        print Y.shape
        del X['Time']
        #Removing mean and scaling to unit variance
        X['Amount'] = StandardScaler().fit_transform(X['Amount'].values.reshape(-1, 1))
        X1 = np.asarray(X)

```

```

y_pred = model.predict(X1)
y_pred = [int(item) for sublist in y_pred for item in sublist]
#print y_pred.shape
print "accuracy", metrics.accuracy_score(Y, y_pred)
print metrics.confusion_matrix(Y, y_pred)
print metrics.classification_report(Y, y_pred)

```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	7680
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 64)	8256
dropout_3 (Dropout)	(None, 64)	0
dense_4 (Dense)	(None, 32)	2080
dropout_4 (Dropout)	(None, 32)	0
dense_5 (Dense)	(None, 1)	33

```

Total params: 50,945
Trainable params: 50,945
Non-trainable params: 0

```

```

None
Train on 1044 samples, validate on 448 samples
Epoch 1/50
0s - loss: 5.1743 - acc: 0.4042 - val_loss: 0.8446 - val_acc: 0.3281
Epoch 2/50
0s - loss: 1.4790 - acc: 0.3314 - val_loss: 1.0919 - val_acc: 0.3281
Epoch 3/50
0s - loss: 1.3967 - acc: 0.3305 - val_loss: 0.9615 - val_acc: 0.3281
Epoch 4/50
0s - loss: 1.1933 - acc: 0.3343 - val_loss: 0.7728 - val_acc: 0.3281
Epoch 5/50
0s - loss: 1.0534 - acc: 0.3879 - val_loss: 0.6780 - val_acc: 0.4799
Epoch 6/50
0s - loss: 0.9248 - acc: 0.5651 - val_loss: 0.6338 - val_acc: 0.7522
Epoch 7/50
0s - loss: 0.8536 - acc: 0.7605 - val_loss: 0.5559 - val_acc: 0.8638

```

Epoch 8/50  
0s - loss: 0.7207 - acc: 0.8439 - val\_loss: 0.4839 - val\_acc: 0.8951  
Epoch 9/50  
0s - loss: 0.6813 - acc: 0.9119 - val\_loss: 0.4420 - val\_acc: 0.8996  
Epoch 10/50  
0s - loss: 0.5412 - acc: 0.9167 - val\_loss: 0.4118 - val\_acc: 0.9062  
Epoch 11/50  
0s - loss: 0.5269 - acc: 0.9234 - val\_loss: 0.4168 - val\_acc: 0.8996  
Epoch 12/50  
0s - loss: 0.4526 - acc: 0.9195 - val\_loss: 0.3912 - val\_acc: 0.8996  
Epoch 13/50  
0s - loss: 0.6253 - acc: 0.9291 - val\_loss: 0.4197 - val\_acc: 0.8906  
Epoch 14/50  
0s - loss: 0.4697 - acc: 0.8841 - val\_loss: 0.4543 - val\_acc: 0.8839  
Epoch 15/50  
0s - loss: 0.4478 - acc: 0.9205 - val\_loss: 0.3316 - val\_acc: 0.9107  
Epoch 16/50  
0s - loss: 0.3840 - acc: 0.9195 - val\_loss: 0.4413 - val\_acc: 0.8772  
Epoch 17/50  
0s - loss: 0.3357 - acc: 0.9176 - val\_loss: 0.3639 - val\_acc: 0.8951  
Epoch 18/50  
0s - loss: 0.3070 - acc: 0.9502 - val\_loss: 0.3096 - val\_acc: 0.9241  
Epoch 19/50  
0s - loss: 0.2654 - acc: 0.9550 - val\_loss: 0.3343 - val\_acc: 0.9241  
Epoch 20/50  
0s - loss: 0.3184 - acc: 0.9511 - val\_loss: 0.3510 - val\_acc: 0.9152  
Epoch 21/50  
0s - loss: 0.2361 - acc: 0.9425 - val\_loss: 0.3329 - val\_acc: 0.9196  
Epoch 22/50  
0s - loss: 0.1889 - acc: 0.9569 - val\_loss: 0.2884 - val\_acc: 0.9353  
Epoch 23/50  
0s - loss: 0.3477 - acc: 0.9502 - val\_loss: 0.4055 - val\_acc: 0.9085  
Epoch 24/50  
0s - loss: 0.2495 - acc: 0.9454 - val\_loss: 0.3276 - val\_acc: 0.9241  
Epoch 25/50  
0s - loss: 0.1852 - acc: 0.9626 - val\_loss: 0.2700 - val\_acc: 0.9397  
Epoch 26/50  
0s - loss: 0.1982 - acc: 0.9674 - val\_loss: 0.2760 - val\_acc: 0.9353  
Epoch 27/50  
0s - loss: 0.2091 - acc: 0.9636 - val\_loss: 0.3116 - val\_acc: 0.9241  
Epoch 28/50  
0s - loss: 0.3697 - acc: 0.9626 - val\_loss: 0.2978 - val\_acc: 0.9308  
Epoch 29/50  
0s - loss: 0.2018 - acc: 0.9521 - val\_loss: 0.3580 - val\_acc: 0.9196  
Epoch 30/50  
0s - loss: 0.2253 - acc: 0.9531 - val\_loss: 0.3407 - val\_acc: 0.9219  
Epoch 31/50  
0s - loss: 0.1935 - acc: 0.9579 - val\_loss: 0.3021 - val\_acc: 0.9286

Epoch 32/50  
0s - loss: 0.2171 - acc: 0.9579 - val\_loss: 0.3012 - val\_acc: 0.9308  
Epoch 33/50  
0s - loss: 0.1752 - acc: 0.9607 - val\_loss: 0.3244 - val\_acc: 0.9241  
Epoch 34/50  
0s - loss: 0.1622 - acc: 0.9607 - val\_loss: 0.3094 - val\_acc: 0.9308  
Epoch 35/50  
0s - loss: 0.1339 - acc: 0.9703 - val\_loss: 0.2647 - val\_acc: 0.9464  
Epoch 36/50  
0s - loss: 0.1047 - acc: 0.9808 - val\_loss: 0.2698 - val\_acc: 0.9487  
Epoch 37/50  
0s - loss: 0.1291 - acc: 0.9703 - val\_loss: 0.2763 - val\_acc: 0.9487  
Epoch 38/50  
0s - loss: 0.0981 - acc: 0.9751 - val\_loss: 0.2710 - val\_acc: 0.9487  
Epoch 39/50  
0s - loss: 0.1477 - acc: 0.9722 - val\_loss: 0.2906 - val\_acc: 0.9442  
Epoch 40/50  
0s - loss: 0.1133 - acc: 0.9674 - val\_loss: 0.3074 - val\_acc: 0.9375  
Epoch 41/50  
0s - loss: 0.1317 - acc: 0.9684 - val\_loss: 0.2840 - val\_acc: 0.9420  
Epoch 42/50  
0s - loss: 0.1229 - acc: 0.9713 - val\_loss: 0.2584 - val\_acc: 0.9487  
Epoch 43/50  
0s - loss: 0.1300 - acc: 0.9741 - val\_loss: 0.2892 - val\_acc: 0.9397  
Epoch 44/50  
0s - loss: 0.1509 - acc: 0.9693 - val\_loss: 0.3375 - val\_acc: 0.9286  
Epoch 45/50  
0s - loss: 0.1307 - acc: 0.9626 - val\_loss: 0.2941 - val\_acc: 0.9397  
Epoch 46/50  
0s - loss: 0.1939 - acc: 0.9665 - val\_loss: 0.2985 - val\_acc: 0.9286  
Epoch 47/50  
0s - loss: 0.1113 - acc: 0.9684 - val\_loss: 0.2752 - val\_acc: 0.9442  
Epoch 48/50  
0s - loss: 0.1086 - acc: 0.9732 - val\_loss: 0.3133 - val\_acc: 0.9397  
Epoch 49/50  
0s - loss: 0.1056 - acc: 0.9665 - val\_loss: 0.2998 - val\_acc: 0.9464  
Epoch 50/50  
0s - loss: 0.1192 - acc: 0.9770 - val\_loss: 0.3017 - val\_acc: 0.9442  
accuracy 0.9419642857142857

[[299 2]

[ 24 123]]

	precision	recall	f1-score	support
0	0.93	0.99	0.96	301
1	0.98	0.84	0.90	147
avg / total	0.94	0.94	0.94	448

```

predicting on whole dataset
(284807, 30)
(284807,)
accuracy 0.9954916838420405
[[283124  1191]
 [    93   399]]

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	284315
1	0.25	0.81	0.38	492
avg / total	1.00	1.00	1.00	284807

### 3.3.5 No Downsampling but provide weights for the classes based on the imbalance

```

In [51]: Y = np.reshape(Y, [Y.shape[0], 1])
         class_weight = {0: 1., 1: 50.}
         X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.3, random_state=42)
         X_train = np.asarray(X_train)
         X_test = np.asarray(X_test)
         Y_train = np.asarray(Y_train)
         Y_test = np.asarray(Y_test)
         K.clear_session()
         InputWidth = X_train.shape[1]
         model = Sequential()
         model.add(Dense(256, input_shape = (InputWidth,), activation='relu'))
         model.add(Dropout(0.2))
         model.add(Dense(128, activation='relu'))
         model.add(Dropout(0.2))
         model.add(Dense(64, activation='relu'))
         model.add(Dropout(0.2))
         model.add(Dense(32, activation='relu'))
         model.add(Dropout(0.2))
         model.add(Dense(1, activation='sigmoid'))
         print model.summary()
         # Compile model
         model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
         output = model.fit(X_train, Y_train, validation_data=(X_test, Y_test), epochs=10, batch_size=128)
         y_pred = model.predict(X_test)
         y_pred = [int(item) for sublist in y_pred for item in sublist]
         #print y_pred.shape
         print "Test accuracy", metrics.accuracy_score(Y_test, y_pred)
         print metrics.confusion_matrix(Y_test, y_pred)
         print metrics.classification_report(Y_test, y_pred)

         y_pred = model.predict(X_train)

```

```

y_pred = [int(item) for sublist in y_pred for item in sublist]
#print y_pred.shape
print "Train accuracy", metrics.accuracy_score(Y_train, y_pred)
print metrics.confusion_matrix(Y_train, y_pred)
print metrics.classification_report(Y_train, y_pred)

```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 256)	7680
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 64)	8256
dropout_3 (Dropout)	(None, 64)	0
dense_4 (Dense)	(None, 32)	2080
dropout_4 (Dropout)	(None, 32)	0
dense_5 (Dense)	(None, 1)	33

```

Total params: 50,945
Trainable params: 50,945
Non-trainable params: 0

```

```

None
Train on 199364 samples, validate on 85443 samples
Epoch 1/10
14s - loss: 0.1443 - acc: 0.9956 - val_loss: 0.0736 - val_acc: 0.9707
Epoch 2/10
12s - loss: 0.0810 - acc: 0.9966 - val_loss: 0.0083 - val_acc: 0.9993
Epoch 3/10
11s - loss: 0.0762 - acc: 0.9974 - val_loss: 0.0159 - val_acc: 0.9987
Epoch 4/10
12s - loss: 0.0864 - acc: 0.9976 - val_loss: 0.0294 - val_acc: 0.9993
Epoch 5/10
11s - loss: 0.0679 - acc: 0.9980 - val_loss: 0.0093 - val_acc: 0.9991
Epoch 6/10
11s - loss: 0.0589 - acc: 0.9969 - val_loss: 0.0169 - val_acc: 0.9979
Epoch 7/10
12s - loss: 0.0608 - acc: 0.9969 - val_loss: 0.0122 - val_acc: 0.9968
Epoch 8/10

```

```

12s - loss: 0.0641 - acc: 0.9968 - val_loss: 0.0264 - val_acc: 0.9972
Epoch 9/10
12s - loss: 0.0584 - acc: 0.9972 - val_loss: 0.0179 - val_acc: 0.9979
Epoch 10/10
13s - loss: 0.0574 - acc: 0.9984 - val_loss: 0.0174 - val_acc: 0.9980
Test accuracy 0.999403110845827
[[85281    26]
 [   25   111]]
      precision    recall  f1-score   support

     0         1.00      1.00      1.00     85307
     1         0.81      0.82      0.81      136

avg / total         1.00      1.00      1.00     85443

Train accuracy 0.9994031018639273
[[198965    43]
 [   76   280]]
      precision    recall  f1-score   support

     0         1.00      1.00      1.00    199008
     1         0.87      0.79      0.82      356

avg / total         1.00      1.00      1.00    199364

```

## Tensorflow + keras

```

In [52]: Y = np.reshape(Y,[Y.shape[0],1])
         X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.33,random_state=42)
         tf.logging.set_verbosity(tf.logging.INFO)
         sess = tf.Session()
         K.set_session(sess)

         InputWidth = X_train.shape[1]
         OutputWidth = 1
         #print Input[0:2]
         #print Label[0:2]
         inp = tf.placeholder(tf.float32, shape=(None,InputWidth))
         out = tf.placeholder(tf.float32, shape=(None,OutputWidth))
         print(out.get_shape())
         print(inp.get_shape())
         predictions=[]
         predictionsTrain =[]
         #OutputWidth = len(y_train[0])
         learning_rate = 0.0001
         #inplayer = Input(shape=(InputWidth, ))

```



```

x = Dense(512,activation='relu')(inp)
print(x)
x = Dense(256,activation='relu')(x)
print(x)
x = Dense(128,activation='relu')(x)
print(x)
#x = K.reshape(x, (len(X_train)*InputWidth,5))
#print x
x = Dense(32,activation='relu')(x)
print(x)
x = Dense(16,activation='relu')(x)
print(x)
preds = Dense(OutputWidth, activation='sigmoid')(x)
print(preds)

loss = tf.reduce_mean(categorical_crossentropy(out, preds))
train_step = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(loss)
#acc_value = accuracy(out, preds)
with sess.as_default():
    sess.run(tf.global_variables_initializer())
    for epoch in range(20):
        print("epoch " + str(epoch))
        #train_step.run(feed_dict={inp:X_train,out:y_train,K.learning_phase(): 1})
        for i in range(1000):
            if(i%1000==0):
                print("iteration number"+str(i))
            _, loss_val = sess.run([train_step, loss],
                                    feed_dict={inp:X_train[i*180:(i+1)*180],out:Y_train[i*180:
print "loss", loss_val
            #train_step.run(feed_dict={inp:X_train[i*180:(i+1)*180],out:X_train[i*180:
#temp=acc_value.eval(feed_dict={inp: X_train,out: y_train})
#print type(temp)
#print temp
#save_path = saver.save(sess, "model_512_5neighbours_ExtraDense.ckpt")
print(out.get_shape())
print(preds.get_shape())
#print preds.eval(feed_dict={inp:X_train[0:2],out: y_train[0:2]})
#p = tf.argmax(preds, axis=1)
predictions = preds.eval(feed_dict={inp: X_test})
predictionsTrain = preds.eval(feed_dict={inp: X_train})
predictions = [int(item) for sublist in predictions for item in sublist]
predictionsTrain = [int(item) for sublist in predictionsTrain for item in sublist]
#print predictions.shape
print Y_test.shape
print np.mean(predictions==Y_test)
print "Train accuracy"
print metrics.confusion_matrix(Y_train, predictionsTrain)
print metrics.classification_report(Y_train, predictionsTrain)

```

```

print "test accuracy"
print metrics.confusion_matrix(Y_test, predictions)
print metrics.classification_report(Y_test, predictions)
#print(len(predictions))
#print(len(predictionsTrain))

```

```

(?, 1)
(?, 29)
Tensor("dense_6/Relu:0", shape=(?, 512), dtype=float32)
Tensor("dense_7/Relu:0", shape=(?, 256), dtype=float32)
Tensor("dense_8/Relu:0", shape=(?, 128), dtype=float32)
Tensor("dense_9/Relu:0", shape=(?, 32), dtype=float32)
Tensor("dense_10/Relu:0", shape=(?, 16), dtype=float32)
Tensor("dense_11/Sigmoid:0", shape=(?, 1), dtype=float32)
epoch 0
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 1
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 2
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 3
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 4
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 5
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 6
iteration number0
loss 0.0

```

```
(?, 1)
(?, 1)
epoch 7
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 8
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 9
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 10
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 11
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 12
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 13
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 14
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 15
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 16
```

```

iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 17
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 18
iteration number0
loss 0.0
(?, 1)
(?, 1)
epoch 19
iteration number0
loss 0.0
(?, 1)
(?, 1)
(93987, 1)
0.0
Train accuracy
[[190477      0]
 [   343      0]]

```

/home/ramchalamkr/.local/lib/python2.7/site-packages/ipykernel\_launcher.py:64: DeprecationWarning

	precision	recall	f1-score	support
0	1.00	1.00	1.00	190477
1	0.00	0.00	0.00	343
avg / total	1.00	1.00	1.00	190820

```

test accuracy
[[93838      0]
 [   149      0]]

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	93838
1	0.00	0.00	0.00	149
avg / total	1.00	1.00	1.00	93987

## 4 Part 2 : Unsupervised Learning approach using Autoencoders

### 5 Steps

#### 5.1 1. Checking and Tidying Data

```
In [9]: X = pd.read_csv('fraud_prep.csv', delimiter=',')
        print X.shape
        del X['Time']
        #Removing mean and scaling to unit variance
        X['Amount'] = StandardScaler().fit_transform(X['Amount'].values.reshape(-1, 1))
```


(284807, 31)

```
In [10]: count_classes = pd.value_counts(X['Class'], sort = True)
        print count_classes
        count_classes.plot(kind = 'bar')
        plt.title("Class distribution")
        plt.xticks(range(2), ['Normal Transaction', 'Fraud Transaction'])
        plt.xlabel("Class")
        plt.ylabel("Frequency")
        plt.show()
```

0 284315

1 492

Name: Class, dtype: int64



```
In [11]: fraud = X[X.Class ==1]
         normal = X[X.Class ==0]
```

```
In [12]: fraud.Amount.describe()
```

```
Out[12]: count    492.000000
         mean      0.135382
         std       1.026242
         min      -0.353229
         25%      -0.349231
         50%      -0.316247
         75%       0.070128
         max       8.146182
         Name: Amount, dtype: float64
```

```
In [13]: normal.Amount.describe()
```

```
Out [13]: count      284315.000000
          mean        -0.000234
          std          0.999942
          min         -0.353229
          25%         -0.330640
          50%         -0.265271
          75%         -0.045177
          max          102.362243
          Name: Amount, dtype: float64
```

```
In [15]: X_train,X_test = train_test_split(X,test_size = 0.3,random_state=42)
          X_train = X_train[X_train.Class==0]
          del X_train['Class']
          y_test = X_test['Class']
          del X_test['Class']
          X_train = np.asarray(X_train)
          X_test = np.asarray(X_test)
          print X_train.shape
          print X_test.shape
          print y_test.shape

(199008, 29)
(85443, 29)
(85443,)
```

## 5.2 2. Model Development

### 5.2.1 Build Model Pure tensorflow

```
In [ ]: inputs = tf.placeholder(tf.float32, shape=(None, X.shape[1]))
          out = X.shape[1]
          def buildmodel(inp):
              e1 = t1.fully_connected(inp, 32, activation_fn=tf.nn.softplus)
              e2 = t1.fully_connected(e1, 16, activation_fn=tf.nn.softplus)
              e3 = t1.fully_connected(e2, 8, activation_fn=tf.nn.softplus)
              d1 = t1.fully_connected(e3, 8, activation_fn=tf.nn.softplus)
              d2 = t1.fully_connected(d1, 16, activation_fn=tf.nn.softplus)
              d3 = t1.fully_connected(d2, out, activation_fn=tf.nn.softplus)
              return out

          outputs = buildmodel(inputs)
          loss = tf.reduce_mean(tf.square(outputs - inputs))
          train_op = tf.train.AdamOptimizer(learning_rate=0.001).minimize(loss)
          #init = tf.global_variables_initializer()
```

### 5.2.2 Keras + Tensorflow

```
In [60]: tf.logging.set_verbosity(tf.logging.INFO)
          sess = tf.Session()
```

```

K.set_session(sess)

InputWidth = X_train.shape[1]
#print Input[0:2]
#print Label[0:2]
inp = tf.placeholder(tf.float32, shape=(None,InputWidth))
out = tf.placeholder(tf.float32, shape=(None,InputWidth))
print(out.get_shape())
print(inp.get_shape())
predictions=[]
predictionsTrain =[]
#OutputWidth = len(y_train[0])
learning_rate = 0.001
#inplayer = Input(shape=(InputWidth, ))
x = Dense(64,activation='relu')(inp)
print(x)
x = Dense(32,activation='relu')(x)
print(x)
x = Dense(16,activation='relu')(x)
print(x)
#x = K.reshape(x, (len(X_train)*InputWidth,5))
#print x
x = Dense(16,activation='relu')(x)
print(x)
x = Dense(32,activation='relu')(x)
print(x)
x = Dense(64,activation='relu')(x)
print(x)
preds = Dense(InputWidth, activation='relu')(x)
print(preds)

loss = tf.reduce_mean(tf.square(preds - inp))
train_step = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(loss)
#acc_value = accuracy(out, preds)
with sess.as_default():
    sess.run(tf.global_variables_initializer())
    for epoch in range(20):
        print("epoch " + str(epoch))
        #train_step.run(feed_dict={inp:X_train,out:y_train,K.learning_phase(): 1})
        for i in range(1000):
            if(i%1000==0):
                print("iteration number"+str(i))
            _, loss_val = sess.run([train_step, loss],
                                   feed_dict={inp:X_train[i*180:(i+1)*180],out:X_train[i*180:
print "loss", loss_val
        #train_step.run(feed_dict={inp:X_train[i*180:(i+1)*180],out:X_train[i*180:
        #temp=acc_value.eval(feed_dict={inp: X_train,out: y_train})
        #print type(temp)

```



```

    #print temp
    #save_path = saver.save(sess, "model_512_5neighbours_ExtraDense.ckpt")
    print(out.get_shape())
    print(preds.get_shape())
    #print preds.eval(feed_dict={inp:X_train[0:2],out: y_train[0:2]})
    #p = tf.argmax(preds, axis=1)
    predictions = preds.eval(feed_dict={inp: X_test})
    print predictions.shape
    mse = np.mean(np.power(X_test - predictions, 2))
    print "test mse",mse
    predictions = preds.eval(feed_dict={inp: X_train})
    mse = np.mean(np.power(X_train - predictions, 2))
    print "train mse",mse
    print(len(predictions))
    print(len(predictionsTrain))

```

```

(?, 29)
(?, 29)
Tensor("dense_18/Relu:0", shape=(?, 64), dtype=float32)
Tensor("dense_19/Relu:0", shape=(?, 32), dtype=float32)
Tensor("dense_20/Relu:0", shape=(?, 16), dtype=float32)
Tensor("dense_21/Relu:0", shape=(?, 16), dtype=float32)
Tensor("dense_22/Relu:0", shape=(?, 32), dtype=float32)
Tensor("dense_23/Relu:0", shape=(?, 29), dtype=float32)
epoch 0
iteration number0
loss 0.43287447
(?, 29)
(?, 29)
epoch 1
iteration number0
loss 0.40982887
(?, 29)
(?, 29)
epoch 2
iteration number0
loss 0.39938778
(?, 29)
(?, 29)
epoch 3
iteration number0
loss 0.3843758
(?, 29)
(?, 29)
epoch 4
iteration number0
loss 0.37673163

```

```
(?, 29)
(?, 29)
epoch 5
iteration number0
loss 0.3734822
(?, 29)
(?, 29)
epoch 6
iteration number0
loss 0.37235656
(?, 29)
(?, 29)
epoch 7
iteration number0
loss 0.37125868
(?, 29)
(?, 29)
epoch 8
iteration number0
loss 0.3702827
(?, 29)
(?, 29)
epoch 9
iteration number0
loss 0.3692525
(?, 29)
(?, 29)
epoch 10
iteration number0
loss 0.36820066
(?, 29)
(?, 29)
epoch 11
iteration number0
loss 0.36791378
(?, 29)
(?, 29)
epoch 12
iteration number0
loss 0.36705187
(?, 29)
(?, 29)
epoch 13
iteration number0
loss 0.3665546
(?, 29)
(?, 29)
epoch 14
```

```

iteration number0
loss 0.36628783
(?, 29)
(?, 29)
epoch 15
iteration number0
loss 0.36620545
(?, 29)
(?, 29)
epoch 16
iteration number0
loss 0.3660761
(?, 29)
(?, 29)
epoch 17
iteration number0
loss 0.36574745
(?, 29)
(?, 29)
epoch 18
iteration number0
loss 0.3655525
(?, 29)
(?, 29)
epoch 19
iteration number0
loss 0.36588037
(?, 29)
(?, 29)
(85443, 29)
test mse 0.6086097037513645
train mse 0.580954868454774
199008
0

```

### 5.2.3 Pure Keras

```

In [17]: InputWidth = X_train.shape[1]
         K.clear_session()
         model = Sequential()
         model.add(Dense(64, input_shape = (InputWidth,), activation='relu'))
         model.add(Dense(32, activation='relu'))
         model.add(Dense(16, activation='relu'))
         model.add(Dense(16, activation='relu'))
         model.add(Dense(32, activation='relu'))
         model.add(Dense(64, activation='relu'))
         model.add(Dense(InputWidth, activation='relu'))

```

```

print model.summary()
# Compile model
model.compile(loss='mean_squared_error', optimizer='adam', metrics=['accuracy'])
op = model.fit(X_train, X_train, validation_data=(X_test, X_test), epochs=30, batch_s

```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 64)	1920
dense_2 (Dense)	(None, 32)	2080
dense_3 (Dense)	(None, 16)	528
dense_4 (Dense)	(None, 16)	272
dense_5 (Dense)	(None, 32)	544
dense_6 (Dense)	(None, 64)	2112
dense_7 (Dense)	(None, 29)	1885

Total params: 9,341  
 Trainable params: 9,341  
 Non-trainable params: 0

None

Train on 199008 samples, validate on 85443 samples

Epoch 1/30  
 10s - loss: 0.7387 - acc: 0.6483 - val\_loss: 0.7038 - val\_acc: 0.7653

Epoch 2/30  
 9s - loss: 0.6716 - acc: 0.7754 - val\_loss: 0.6895 - val\_acc: 0.7953

Epoch 3/30  
 9s - loss: 0.6622 - acc: 0.7953 - val\_loss: 0.6840 - val\_acc: 0.8092

Epoch 4/30  
 10s - loss: 0.6572 - acc: 0.8069 - val\_loss: 0.6809 - val\_acc: 0.8201

Epoch 5/30  
 10s - loss: 0.6568 - acc: 0.8062 - val\_loss: 0.6791 - val\_acc: 0.8167

Epoch 6/30  
 11s - loss: 0.6530 - acc: 0.8155 - val\_loss: 0.6774 - val\_acc: 0.8266

Epoch 7/30  
 9s - loss: 0.6518 - acc: 0.8198 - val\_loss: 0.6797 - val\_acc: 0.7824

Epoch 8/30  
 10s - loss: 0.6521 - acc: 0.8185 - val\_loss: 0.6772 - val\_acc: 0.8182

Epoch 9/30  
 9s - loss: 0.6512 - acc: 0.8227 - val\_loss: 0.6759 - val\_acc: 0.8289

Epoch 10/30  
 9s - loss: 0.6511 - acc: 0.8237 - val\_loss: 0.6731 - val\_acc: 0.8175

```

Epoch 11/30
9s - loss: 0.6459 - acc: 0.8283 - val_loss: 0.6703 - val_acc: 0.8396
Epoch 12/30
9s - loss: 0.6448 - acc: 0.8333 - val_loss: 0.6703 - val_acc: 0.8402
Epoch 13/30
9s - loss: 0.6443 - acc: 0.8350 - val_loss: 0.6698 - val_acc: 0.8429
Epoch 14/30
9s - loss: 0.6444 - acc: 0.8345 - val_loss: 0.6692 - val_acc: 0.8448
Epoch 15/30
9s - loss: 0.6438 - acc: 0.8362 - val_loss: 0.6736 - val_acc: 0.7922
Epoch 16/30
9s - loss: 0.6436 - acc: 0.8359 - val_loss: 0.6687 - val_acc: 0.8427
Epoch 17/30
9s - loss: 0.6429 - acc: 0.8385 - val_loss: 0.6719 - val_acc: 0.8133
Epoch 18/30
9s - loss: 0.6430 - acc: 0.8379 - val_loss: 0.6683 - val_acc: 0.8466
Epoch 19/30
9s - loss: 0.6430 - acc: 0.8384 - val_loss: 0.6679 - val_acc: 0.8425
Epoch 20/30
9s - loss: 0.6426 - acc: 0.8393 - val_loss: 0.6681 - val_acc: 0.8431
Epoch 21/30
9s - loss: 0.6380 - acc: 0.8450 - val_loss: 0.6602 - val_acc: 0.8450
Epoch 22/30
9s - loss: 0.6334 - acc: 0.8485 - val_loss: 0.6588 - val_acc: 0.8507
Epoch 23/30
9s - loss: 0.6309 - acc: 0.8557 - val_loss: 0.6560 - val_acc: 0.8619
Epoch 24/30
9s - loss: 0.6305 - acc: 0.8568 - val_loss: 0.6558 - val_acc: 0.8626
Epoch 25/30
9s - loss: 0.6290 - acc: 0.8599 - val_loss: 0.6549 - val_acc: 0.8584
Epoch 26/30
9s - loss: 0.6296 - acc: 0.8574 - val_loss: 0.6548 - val_acc: 0.8636
Epoch 27/30
9s - loss: 0.6291 - acc: 0.8586 - val_loss: 0.6550 - val_acc: 0.8620
Epoch 28/30
9s - loss: 0.6313 - acc: 0.8530 - val_loss: 0.6537 - val_acc: 0.8674
Epoch 29/30
9s - loss: 0.6278 - acc: 0.8601 - val_loss: 0.6564 - val_acc: 0.8550
Epoch 30/30
9s - loss: 0.6270 - acc: 0.8551 - val_loss: 0.6456 - val_acc: 0.8690

```


### 5.3 3. Model Performance

```

In [18]: plt.plot(op['loss'],label='train')
         plt.plot(op['val_loss'],label='test')
         plt.legend(bbox_to_anchor=(0., 1.02, 1., .102), loc=3,
                   ncol=2, mode="expand", borderaxespad=0.)

```


```
plt.title('Loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.show()
```




```
In [19]: y_pred = model.predict(X_test)
mse = np.mean(np.power(X_test - y_pred, 2), axis=1)

In [23]: reconstructionLoss = pd.DataFrame({'ActualClass':y_test, 'mse':mse})
Normal = reconstructionLoss[(reconstructionLoss['ActualClass'] == 0) & (reconstructionLoss['mse'] < 0.05)]
Fraud = reconstructionLoss[(reconstructionLoss['ActualClass'] == 1)]

In [24]: plt.hist(Normal['mse'],bins=10)
plt.title("reconstruction loss for normal transaction")
plt.xlabel("mse")
plt.ylabel("Frequency")
plt.show()
```



```
In [25]: plt.hist(Fraud['mse'],bins=10)
plt.title("reconstruction loss for fraud transaction")
plt.xlabel("mse")
plt.ylabel("Frequency")
plt.show()
```



```
In [26]: def TN_TP(ytrue,ypred,l,tpe):
    count = 0
    for i in range(l):
        if(ytrue[i] == tpe and ypred[i] == tpe):
            count+=1
    return count

def FN_FP(ytrue,ypred,l,tpe):
    count = 0
    for i in range(l):
        if(ytrue[i] == tpe and ypred[i] != tpe):
            count+=1
    return count

def calculate_TPR_FPR(ytrue,ypred,P):
    TP = TN_TP(ytrue,ypred,len(P),1)
    TN = TN_TP(ytrue,ypred,len(P),0)
    FP = FN_FP(ytrue,ypred,len(P),0)
    FN = FN_FP(ytrue,ypred,len(P),1)

    TPR = float(TP)/(TP+FN)
    FPR = float(FP)/(FP+TN)

    return TPR,FPR
```



```

def plot_roc(TPRFinal,FPRFinal):
    roc = auc(FPRFinal,TPRFinal)
    plt.figure()
    lw = 2
    plt.plot(FPRFinal, TPRFinal, color='darkorange',lw=lw,label='ROC curve (area = %0
    plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
    plt.xlim([-0.1, 1.1])
    plt.ylim([-0.1, 1.1])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver operating characteristic example')
    plt.legend(loc="lower right")
    plt.show()

```


```

In [27]: TPRFinal = []
        FPRFinal = []
        Treshold = [0.5,1,1.5,2,2.5,3,3.5,4]
        for i in Treshold:
            y_pred = [1 if x>i else 0 for x in reconstructionLoss['mse']]
            TPR,FPR = calculate_TPR_FPR(reconstructionLoss['ActualClass'].values,y_pred,y_pre
            TPRFinal.append(TPR)
            FPRFinal.append(FPR)

        #print TPRFinal
        #print FPRFinal

        plot_roc(TPRFinal,FPRFinal)

```



```
In [28]: Treshold = [0.5,1,1.5,2,2.5,3,3.5,4,4.5,5,5.5]
        for i in Treshold:
            y_pred = [1 if x>i else 0 for x in reconstructionLoss['mse']]
            print metrics.confusion_matrix(reconstructionLoss['ActualClass'], y_pred)
            print metrics.classification_report(reconstructionLoss['ActualClass'], y_pred)
```

```
[[60575 24732]
 [   6  130]]
      precision    recall  f1-score   support

     0       1.00      0.71      0.83     85307
     1       0.01      0.96      0.01        136

avg / total       1.00      0.71      0.83     85443
```

```
[[78740 6567]
 [  13   123]]
      precision    recall  f1-score   support

     0       1.00      0.92      0.96     85307
     1       0.02      0.90      0.04        136

avg / total       1.00      0.92      0.96     85443
```

```
[[81722 3585]
 [  15   121]]
      precision    recall  f1-score   support

     0       1.00      0.96      0.98     85307
     1       0.03      0.89      0.06        136

avg / total       1.00      0.96      0.98     85443
```

```
[[82783 2524]
 [  20   116]]
      precision    recall  f1-score   support

     0       1.00      0.97      0.98     85307
     1       0.04      0.85      0.08        136

avg / total       1.00      0.97      0.98     85443
```

```
[[83356 1951]
 [  24   112]]
      precision    recall  f1-score   support
```

0	1.00	0.98	0.99	85307
1	0.05	0.82	0.10	136

avg / total	1.00	0.98	0.99	85443
-------------	------	------	------	-------

[[83695 1612]  
[ 26 110]]

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	0.98	0.99	85307
1	0.06	0.81	0.12	136

avg / total	1.00	0.98	0.99	85443
-------------	------	------	------	-------

[[83943 1364]  
[ 31 105]]

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	0.98	0.99	85307
1	0.07	0.77	0.13	136

avg / total	1.00	0.98	0.99	85443
-------------	------	------	------	-------

[[84157 1150]  
[ 36 100]]

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	0.99	0.99	85307
1	0.08	0.74	0.14	136

avg / total	1.00	0.99	0.99	85443
-------------	------	------	------	-------

[[84317 990]  
[ 42 94]]

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	0.99	0.99	85307
1	0.09	0.69	0.15	136

avg / total	1.00	0.99	0.99	85443
-------------	------	------	------	-------

[[84440 867]  
[ 44 92]]

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	0.99	0.99	85307
1	0.10	0.68	0.17	136

avg / total	1.00	0.99	0.99	85443
-------------	------	------	------	-------

```
[[84536  771]
 [   47   89]]
```

	precision	recall	f1-score	support
0	1.00	0.99	1.00	85307
1	0.10	0.65	0.18	136

avg / total	1.00	0.99	0.99	85443
-------------	------	------	------	-------